Congratulations! You passed!

Grade received 100% **Latest Submission Grade** 100% **To pass** 80% or higher

Retake the assignment in **23h 54m**

Go to next item

1. True/False: Suppose you learn a word embedding for a vocabulary of 60000 words. Then the embedding vectors could be 60000 dimensional, so as to capture the full range of variation and meaning in those words. **1 / 1 point**

- ○ True
- ◉ False

Expand

Correct
No, the dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 1000.

2. True/False: t-SNE is a linear transformation that allows us to solve analogies on word vectors. **1 / 1 point**

- ◉ False
- ○ True

Expand

Correct
tr-SNE is a non-linear dimensionality reduction technique.

3. Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set. **1 / 1 point**

x (input text)	y (happy?)
Having a great time!	1
I'm sad it's raining.	0
I'm feeling awesome!	1

Even if the word "wonderful" does not appear in your small training set, what label might be reasonably expected for the input text "I feel wonderful!"?

- ◉ y=1
- ○ y=0

Expand

Correct

Yes, word vectors empower your model with an incredible ability to generalize. The vector for “wonderful” would contain a negative/unhappy connotation which will probably make your model classify the sentence as a "1".

4. Which of these equations do you think should hold for a good word embedding? (Check all that apply) **1 / 1 point**

- [x] $e_{boy} - e_{brother} \approx e_{girl} - e_{sister}$

  ✓ **Correct**
  Yes!

- [ ] $e_{boy} - e_{brother} \approx e_{sister} - e_{girl}$
- [x] $e_{boy} - e_{girl} \approx e_{brother} - e_{sister}$

  ✓ **Correct**
  Yes!

- [ ] $e_{boy} - e_{girl} \approx e_{sister} - e_{brother}$

Expand

✓ **Correct**
Great, you got all the right answers.

5. Let A be an embedding matrix, and let o_{4567} be a one-hot vector corresponding to word 4567. Then to get the embedding of word 4567, why don’t we call $A * o_{4567}$ in Python? **1 / 1 point**

- ( ) The correct formula is $A^T * o_{4567}$
- ( ) This doesn’t handle unknown words (<UNK>).
- (•) It is computationally wasteful.
- ( ) None of the answers are correct: calling the Python snippet as described above is fine.

Expand

✓ **Correct**
Yes, the element-wise multiplication will be extremely inefficient.

6. When learning word embeddings, we create an artificial task of estimating $P(target \mid context)$. It is okay if we do poorly on this artificial prediction task; the more important by-product of this task is that we learn a useful set of word embeddings. **1 / 1 point**

- (•) True
- ( ) False

Expand

✓ **Correct**

7. In the word2vec algorithm, you estimate $P(t \mid c)$, where t is the target word and c is a context word. How are t and c chosen from the training set? Pick the best answer. **1 / 1 point**

- (●) c and t are chosen to be nearby words.
- ( ) c is the one word that comes immediately before t
- ( ) c is a sequence of several words immediately before t
- ( ) c is the sequence of all the words in the sentence before t

Expand

✓ **Correct**

8. Suppose you have a 10000 word vocabulary, and are learning 100-dimensional word embeddings. The word2vec model uses the following softmax function: **1 / 1 point**

$$P(t \mid c) = \frac{e^{\theta_t^T e_c}}{\sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c}}$$

True/False: After training, we should expect θ_t to be very close to e_c when t and c are the same word.

- ( ) True
- (●) False

Expand

✓ **Correct**
To review this concept watch the *Word2Vec* lecture.

9. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective: **1 / 1 point**

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - log X_{ij})^2$$

Which of these statements are correct? Check all that apply.

- [x] Theoretically, the weighting function $f(.)$ must satisfy $f(0) = 0$

✓ **Correct**

- [x] θ_i and e_j should be initialized randomly at the beginning of training.

✓ **Correct**

- [ ] θ_i and e_j should be initialized to 0 at the beginning of training.
- [x] X_{ij} is the number of times word j appears in the context of word i.

✓ **Correct**

Expand

Correct
Great, you got all the right answers.

10. You have trained word embeddings using a text dataset of t_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of t_2 words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful? **1 / 1 point**

- When t_1 is smaller than t_2
- When t_1 is larger than t_2
- When t_1 is equal to t_2

Expand

Correct
Transfer embeddings to new tasks with smaller training sets.